Chart-to-Table Conversion: A Survey

Mohamed Fayed, Kruthik Ravikanti, Bruce Walker

mohamed.fayed@gatech.edu, kravikanti3@gatech.edu,

bruce.walker@psych.gatech.edu

Abstract—Multimodal Large Language Models (MLLMs) have demonstrated impressive visual capabilities in many Visual Question Answering (VQA) tasks. Chart-to-Table Conversion is the problem of extracting structured tabular data out of input images, enabling better accessibility and higher storage efficiency. In this survey, we review the recent advancements in Chart-to-Table task, analyze existing datasets and benchmark some state-of-the-art models. Our quantitative and qualitative analysis highlights strengths and limitations of current approaches in handling complex charts, preserving numerical precision and generalizing across various types of charts. Our study reveals significant room for improvement in Chart-to-Table Conversion, emphasizing the need for more specialized and robust methods.

CONTENTS

1 INTRODUCTION

Data visualization is an essential tool for communicating information efficiently, with charts and graphs playing a fundamental role in various domains, includ-

ing finance, healthcare, and scientific research. These visualizations are usually saved in form of images, and thus require extracting structured tabular data from chart images remains a challenging problem, particularly for automated systems. This challenge is especially critical for accessibility, where individuals with visual impairments rely on machine-readable formats to interpret visual information.

Chart-to-Table is the task of extracting data points from an image of a chart into a table usually in markdownLiu et al., Masry et al., 2022a, 2024b. It was usually tackled via combination of rule-based tecniques, image processing, CNN-based deep learning architectures and optical character recognition (OCR). These techniques struggle with diverse chart formats, occlusions and variations in visual styles. More recently, this task was tackled via transformer-based models and Multimodal Large Language Models (MLLMs), which provide more robust solutions.

In this survey, we provide a comprehensive review of Chart-to-Table conversion techniques, covering benchmark datasets, model architectures, evaluation metrics, and existing challenges. Furthermore, we assess the performance of several state-of-the-art transformers-based models, including MLLMs, on this task, highlighting their strengths and weaknesses. Our findings suggest that while these models exhibit promising results, significant challenges remain, including generalization across different chart types, handling visual input effectively, and improving numerical fidelity.

Our main contributions are:

1. Survey recent advancements in Chart-to-Table task; to the best of our knowledge, this is the first survey focusing on Chart-to-Table Conversion,
2. Conduct a comprehensive qualitative and quantitative analysis on various benchmarks covering various common types of charts, and
3. highlight strengths and limitations of those models, emphasizing on rooms for improvement in tackling Chart-to-Table Conversion.

2 CHART-TO-TABLE IN LITERATURE

There has been a lot of work related to chart analysis in multiple dimensions. Some researchers focused on converting charts to tables (Chart-to-Table Conversion). Others concentrated on Chart Question Answering and Chart Summarization. In this work, we focus on Chart-to-Table Conversion, which requires tackling Optical Character Recognition (OCR) and comprehensive Visual Understanding. In this

section, we review related work in Traditional Computer Vision, Deep Learning and Multimodal Large Language Models and how they influenced the evolution of this task.

2.1 Computer Vision and Deep Learning for Chart-to-Table Conversion

It all begins with computer vision (CV)-based methods. Early attempts to extract tabular data from charts utilized traditional image processing techniques like segmentation and edge detection that could identify visual elements like bars, axes, and legends Sreevalsan-Nair, Dadhich, and Daggubati, 2021. These approaches were effective to some extent, but they struggled with the many variations found in chart styles and occlusions. Furthermore, rule-based approaches were often unsuccessful due to complex charts that had overlapping elements or non-standard layouts Poco and Heer, 2017. As a result, researchers pursued more robust techniques that could generalize better across diverse chart designs.

To deal with generalization problem, manny efforts were put towards integrating deep learning solutions. Systems such as ChartSense Jung et al., 2017 and ChartOCR Luo et al., 2021 utilized convolutional neural networks (CNNs) to classify types of charts and extract available data with higher accuracies. For instance, ChartOCR combined rule-based methods with deep CNNs to balance generalization and accuracy, achieving strong performance across multiple chart types Luo et al., 2021. However, these methods needed extensive labeled datasets and lacked generalization across different chart formats.

There were many efforts towards handling specific types of charts using either Computer Vision or Deep Learning tecniques. ChartParser Kumar, Ganu, and Guha, 2022 was a pipeline that focuses on extracting data points for various kinds of bar charts. They build rule based algorithm using OCR and image processing techniques. Another group of researchers Li et al., 2022 focused on extracting underlying data points for camera photos of line charts using custom architecture based on Faster-RCNN Ren et al., 2015. They gathered synthetic and real datasets for line charts to improve the performance of their model. Their approach shows promising direction in utilizing synthetic data in improving models on real world test images. This direction of research highlights challenges of each kind of chart, but it has many disconnected componenets and their results are not necessarily comparable.

4

2.2 Transformer Models

After the introduction of transformers Vaswani et al., 2017, Pretraining and Fine-tuning paradigm caught researchers attention. Masry et al., 2023 gathered a large dataset of 6.9M questions and charts, and used it to pre-train UniChart model. Another group of researchers Liu et al., 2022b introduced Matcha model, which was a fine-tuned version of Pix2Struct Lee et al., 2023 on many tasks including Chart-to-Table. Later on, Liu et al., 2022a continued pre-training it on Chart-to-Table only to create DePlot. They forwarded the generated table, human query and an example to FlanPaLM 540B Chung et al., 2024 to answer complex queries. Cheng, Dai, and Hauptmann, 2023 took a different approach by training a transformers based chart component detection and combine it with extended pretrained T5 Raffel et al., 2020 or TaPas Herzig et al., 2020 models.

2.3 Multimodal Large Language Models (MLLMs)

MLLMs have made significant progress in many tasks and Chart-to-Table was no exception. One direction is to utilize general domain LLMs without any tuning. This direction includes Prompting and Retrieval Augmented Generation techniques Cao et al., Voigt et al., 2024, 2023, which improve the capabilities of general LLMs on chart related tasks.

Another direction is to fine-tune LLMs on Chart-specific Instruction following datasets. Masry et al., 2024a introduced an instruction-following dataset and instruction tuned both LLaMA 2 7B Touvron et al., 2023 and Flan-T5 XL 3B Chung et al., 2024 on it. This instruction tuning strategy proved to make a generalized model that can handle unseen chart related tasks. Similarly, Masry et al., 2024b fine-tuned PaliGemma Beyer et al., 2024 to create a 3B ChartGemma. Another key distinguishing contribution was the method of instruction-following data generation. They generated for predefined tasks, such as Chain of Thought and Chart-to-Tables in form of markdown, and open-ended tasks, such as justifying trends in charts and describing visual elements.

A unique direction is chart-to-code generation, which aims to provide lossless representations of charts. Matcha Liu et al., 2022b was the first model to include Chart-to-Code within its training tasks. ChartCoder Zhao et al., 2025 is a new model that translates chart images into executable Python code (e.g., using matplotlib). Trained on the newly released Chart2Code-160k dataset containing 160k

chart-code pairs across 27 chart types, ChartCoder uses a Code LLM backbone and a Snippet-of-Thought prompting method for step-by-step generation. While intended for re-rendering, the output code preserves all chart data, making it a viable route for Chart-to-Table extraction through code interpretation.

Another recent model, ChartCitor Goswami et al., 2025, focuses on improving how chart question-answering models support their answers with evidence. Instead of just generating answers. ChartCitor uses a group of LLM agents to first extract a table from a chart, then break the answer into smaller claims, and finally find the exact table cells that support those claims. It also connects these cells back to specific parts of the chart using object detection and visual prompts. ChartCitor does not aim to extract full tables, it introduces a useful method for linking visual data to structured output and improves transparency in chart understanding tasks, which makes it a good candidate for verifying correctness of generated tables.

3 EXPERIMENTS

3.1 Datasets

Many datasets were introduced for various chart related tasks, with no special attention to Chart-to-Table Conversion. After investigation, we found that some datasets include chart images with their corresponding data points. In our analysis, we focus on reporting scores on two of those testsets: ICPR22 Rousseau and Kapralos, 2023 and PlotQA Methani et al., 2020.

ICPR22 testset Rousseau and Kapralos, 2023—was scrapped from research papers published on PubMed Central website.[1] Those publications are in biomedical and life sciences domains. It contains 443 charts splitted into 5 types: Line Charts, Horizontal and Vertical Bar Charts, Scatter Plot and Vertical Box Plot. It represents a source of real-world generated charts, as they were generated by a large set of researchers.

PlotQA Methani et al., 2020—was made by scrapping data from various online sources, such as World Bank and Open Data, generate plots out of these data points, and ask annotators questions about those provided charts. In our work, we focus on the data points used in constructing the charts only. Its test set contains 33657 charts divided equally among dotted line charts, line charts, and vertical

---

1 https://pmc.ncbi.nlm.nih.gov

and horizontal bar charts. We hypothesized that there is no need to test on all 33k charts, so we performed comparison among scores for many sample sizes. Table 1 shows that it is enough to compute scores on a sample of 4k randomly selected charts.

Test-Sample:—For the sake of qualitative analysis, we selected a set of charts from PlotQA and ICPR22 testsets. We randomly selected from each kind of chart in PlotQA. However, for ICPR22, we selected charts from various data sizes, i.e. of various sizes of json files. We noticed that larger file sizes are directly correlated to larger count of data points of corresponding charts.

3.2 Models

For our study, we tested the following models in zero-shot setting:

- Gemini 1.5 Flash Team et al., 2024: A general purpose lightweight MLLM, which is the largest in this list.
- ChartGemma Masry et al., 2024b: A specialized model in chart summarization, question answering and reasoning about charts. It utilizes PaliGemma 4B Beyer et al., 2024 as its backbone, and was tuned on Visual Chart Instructions dataset.
- DePlot Liu et al., 2022a: A 282M parameters Vision Encoder-Text Decoder transformer architecture. It was trained on Chart-to-Table Conversion task.

3.3 Prompts

For DePlot, we used the prompt recommended in its Hugging Face Model Card: "Generate underlying data table of the figure below:". For ChartGemma, we used a simple prompt of "program of thought: extract data points and formulate them into markdown table. Note: all data points should be either integer or float. Do not write percentages as '56

For Gemini 1.5 Flash, we used a longer prompt, and if its output is not parse-able into json object, we prompt it again to fix its format. The main prompt is:

```
"System Message:    You are an EXPERT in scientific visualizations analysis and ch
You should help people with disabilities in giving an overview of the chart and con
    prompt: For the following image, tell me the name of x and y axes, and the valu
    The output should be in format of
    {x_axis_name: x_axis_vavlues_list, y_xis_name: y_axis_values_list}"
```

3.4 Evaluation

There has been many metrics for evaluating Chart-to-Table conversion. In Luo et al., 2021, they introduced different metric for each kind of chart:

1. For Bar Charts, defined a custom distance function for pairwise point comparison and find minimum cost between prediction and ground truth,
2. For Line charts, it is handled as continuous similarity problem. For each predicted line, it computes the pointwise error between it and each ground truth line, and choose the minimal value.
3. For Pie Charts, they consider its scoring as sequence matching problem, thus solved using dynamic programming.

Relative Number Similarity Score Masry et al., 2022, also known as Relaxed Accuracy Measure, is to compute highest accuracy of generated numbers relative to ground truth. However, this evaluation approach has two main drawbacks:

1. It does not consider the position of numbers within the table, and
2. It ignores textual errors.

To overcome those limitations, Relative Mapping Similarity (RMS) Liu et al., 2022a computes edit distance between columns names, compute accuracy of values within columns of least edit distances and compute F1-score. [2] To learn more about the algorithm, please check appendix A

To compare among those models, we compute Relative Mapping Similarity (RMS) Liu et al., 2022a and do qualitative analysis on our Test-Sample testset.

Success rate:—While testing, we noticed that all models suffer in generating markdowns/jsons that can be correctly parsed by pandas with various levels. To quantify this behavior, we introduce Success Rate (SR) for each model. Success Rate (SR) is the count of correctly parsed charts ($C_{correct}$) divided by count of predicted charts within the testset by selected model (C_{pred}) 1.

$$SR = C_{correct}/C_{pred} \tag{1}$$

---

2 our RMS implementation can be found here.

4 RESULTS AND DISCUSSION

4.1 Quantitative Analysis

4.1.1 *Scores*

Table 2 shows the final scores for all models on PlotQA and ICPR22 testsets. We were able to generate around 7000 charts from PlotQA. We hypothesized that the size of 33k test images is very large and we should not find significant difference between scores on 3k and 33k. To validate this hypothesis, we computed scores for 3k, 4k, 5k and 6k charts to see whether there is a significant need to infere on all 33k images. Additionally, we computed the scores on the entire testset using DePlot model. As shown in table 1, there are small differences across samples of 4k, 5k and k relative to variations among scores of all models. This means that a sample of 4k is sufficient to provide a reliable estimate to model's performance.

	Gemini 1.5 Flash	
Size	SR	RMS
1.5k	86.41%	55.43%
3k	85.94%	55.17%
4k	8606%	55.26%
5k	86.39%	55.56%
6k	86.40%	55.56%

	ChartGemma	
1.5k	15.93%	24.93%
3k	17.53%	29.66%
4k	17.21%	30.26%
5k	18.15%	31.26%
6k	18.32%	31.17%
6946	18.94%	31.24%

	Deplot	
Size	SR	RMS
1.5k	83.6%	83.63%
3k	84.57%	83.72%
4k	84.65%	83.77%
5k	84.24%	83.62%
6k	83.93%	83.74%
6946	83.7%	83.66%
33k (all)	83.41%	83.68%

Table 1—Scores of all models for different sample sizes from PlotQA. The first 2 digits are the same starting at sample size of 4k, which implies we do not need more than 4k images to get scores of precision of 2 digits.

Model	PlotQA		ICPR22	
	SR	RMS	SR	RMS
Gemini 1.5 Flash	86.5%	55.6%	76.8%	19%
ChartGemma	18.9%	31.2%	48.4%	17.2%
Deplot	83.7%	83.66%	75.8%	19.6%

Table 2—Models scores on 7k subset of PlotQA and all ICPR22 testsets.

Fianl scores are shown in table 2. Gemini 1.5 Flash achieved the highest Success Rate (SR), indicating its strong ability to produce well-formatted and parse-able outputs. However, DePlot outperformed all models in terms of Relative Mapping Similarity (RMS), showcasing its superior alignment with the ground truth table values. On the other hand, ChartGemma recorded the lowest performance in both SR and RMS.

The relatively high SR of Gemini 1.5 Flash suggests that a powerful text decoder plays a crucial role in generating structured and clean outputs. Meanwhile, De-Plot's strong RMS score—despite being a smaller, specialized model—underscores the value of task-specific architectural tuning. These results suggest that instruction fine-tuning (as in ChartGemma) may not be sufficient to achieve high-quality Chart-to-Table conversions without deeper task alignment and robust visual reasoning.

The low scores of all models in ICPR22 testset indicate that it is challenging testset. Upon manual inspection, we identified many contributors to low RMS scores. One major issue is the absence of labels on many line charts. Many columns in tables of the testset contain "unnamed data series 0", which has 1 normalized edit distance with whatever models generate as columns names, thus 0 RMS score occurs frequently. In a future work, we would also take into account computing RMS both with and without considering normalized edit distance. Another reason for those low scores is the failure of all models in properly extracting data from scatter plots and vertical box plots, as illustrated in section B.1 and section B.2 respectively.

4.2 Qualitative Analysis

4.2.1 *Text Recognition*

As per the analyzed sample, we found a single error in DePlot's output in recognizing name of second row in table 28308 in PlotQA. Other than that, all text was correctly recognized by all models. However, tables 6, 5 and 4 show models' tendency to labelize even if there are no labels in the input image similar to figure 1.. 3

4.2.2 *Values Extraction*

For PlotQA and ICPR22 samples, it is frequent to find errors like:

1. rounding errors, e.g. $15.42-> 15$ and $15.6-> 15$.
2. Precision Errors: we have noticed that the models can not predict more than 3 digits for Gemini 1.5 Flash and ChartGemma or 4 digits for DePlot per each value, e.g. $126765000.0-> 156000000$.
3. In case of near values, e.g. $24.18, 24.09$, there might be some errors, e.g. predicting 23 instead of 24. That kind of error may result in changing trend, e.g. steady performance may seam as decreasing. 4
4. All models can differentiate outputs based on scale, e.g. 156000000&50.2 for instance. all models sometimes mispredict scale, e.g. table 5 where ChartGemma returned values multiplied by 10 relative to ground truth. However, DePlot was found to be the most robust against this error.
5. Occasionally, ChartGemma and Gemini 1.5 Flash swap two columns as shown in table 11. As a result, RMS score is significantly lower (f1=34%) than its fixed version (f1=83%).

In the remaining part of this section, we illustrate issues related to each kind of graphs.

Line Charts—

1. Tables 6, 5 and 4 show that both models are very good in extracting data points from line charts.

---

3 In some cases, the ground truth is mislabelled. The reference has no values for x-axis, but the image includes them as in 3.

4 It is worth noting that we have not seen cases where increasing is replaced by decreasing trends or vice versa by any model in our Test-Sample.

Figure 1—Example for Line Chart from PlotQA testset # 21673 about Portfolio Investment

	name	color	label	bboxes
0	Portfolio Investment	#BA55D3	Portfolio Investment	['y': 51, 'x': 132, 'w': 466, 'h': 413, 'y': 51, '

Table 3—Reference for Line Chart from PlotQA #21673 Portfolio Investment

	Year	Investment (in USD)
0	2008	54000000
1	2009	240000000
2	2010	200000000

Table 4—Predicted data points by Gemini 1.5 Flash for Line Chart from PlotQA #21673 Portfolio Investment

	Year	Investment (in USD)
1	2008	500000000
2	2009	2400000000
3	2010	1900000000

Table 5—ChartGemma: prediction for PlotQA line chart #21673

Year	Portfolio Investment
2008	55000000
2009	241000000
2010	192700000

Table 6—Predicted data points by Deplot for Line Chart from PlotQA #21673 Portfolio Investment

1. There are some graphs, like 2, the Gemini API just fails with no clear response message. However, it is suspected that the large number of data points might be the reason. It is noticable for the same chart to cause DePlot to repeat some rows as shown in table 7.


Figure 2—Example for line charts that causes the API to fail.

1. Table 10 shows that ChartGemma may fail in extracting data points from slightly complex graphs. It fails in both extracting correct values as well as mapping them to the correct label.

Bar Charts—

1. All models are good in extracting data points of Horizontal and Vertical bar charts. Tables 12 and 13 include Gemini 1.5 Flash and ChartGemma predictions for figure 5 respectively. However, DePlot is better in precision and in extracting negative points as in Figure 6.

Other Types of Charts:—ICPR22 testset includes scatter and vertical box plots. However, the performance of all models is very poor in extracting underlying data points in both types. Appendix **??** include examples for predictions of all models.

5 CONCLUSION AND RECOMMENDATIONS

In conclusion, our study provides a comprehensive evaluation of the performance of various models on chart-to-table Conversion. Our results show that while some models excel in certain types of charts, none of them perform evenly across all

Wavelength (nm)	C₄	TNT	AN	Soil	Vegetation
	23.7	26.8	29.6	23.7	29.4
	23.9	32.6	35.6	27.9	33.8
70.0	23.9	28.6	46.9	25.6	30.1
50.0	25.5	33.2	36.5	29.9	33.8
50.0	25.2	35.6	48.3	29.5	25.2
50.0	25.6	34.8	48.3	30.9	25.4
50.0	25.5	34.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2
50.0	25.8	33.8	48.3	30.5	25.2

Table 7—Deplot prediction for ICPR22 line chart #6339

Figure 3—PlotQA # 20049: Line chart containing 4 lines.

	name	color	label	bboxes
0	Argentina	#228B22	Argentina	['y': 386, 'x': 101, 'w': 320, 'h': 58, 'y': 386, 'x': 421, 'w':
1	Mauritania	#CD853F	Mauritania	['y': 186, 'x': 101, 'w': 320, 'h': 4, 'y': 150, 'x': 421, 'w':
2	New Zealand	#66CDAA	New Zealand	['y': 368, 'x': 101, 'w': 320, 'h': 23, 'y': 391, 'x': 421, 'w':
3	Togo	#90EE90	Togo	['y': 60, 'x': 101, 'w': 320, 'h': 38, 'y': 56, 'x': 421, 'w': 32

Table 8—Reference table for PlotQA line chart # 20049

	Year	Argentina	Mauritania	New Zealand	Togo
0	2000	10.600000	26.000000	15.800000	31.400000
1	2003	14.200000	25.800000	14.000000	33.200000
2	2004	13.600000	28.000000	13.000000	33.800000
3	2005	13.400000	30.200000	14.600000	30.000000

Table 9—Gemini 1.5 Flash prediction for PlotQA line chart # 20049

	Years	Argentina	Mauritius	New Zealand	Togo
1	2000	21	15	22	11
2	2003	21	14	23	14
3	2004	22	13	23	13
4	2005	21	14	21	14

Table 10—ChartGemma prediction for PlotQA line chart # 20049. The model fails in mapping lines with values, e.g. Togo column seams more likely to be Argentina. For values, it is obvious that ChartGemma is very far away from correctly detecting values greater than 20!

Figure 4—A good example for graph in the wild that causes Gemini 1.5 Flash to fail.

	Australia	Turkmenistan
0	'Year': 2009.0, 'Subscribers per 100 People': 47.0	'Year': 2009.0, 'Subscribers per 100 People': 48.5
1	'Year': 2010.0, 'Subscribers per 100 People': 46.0	'Year': 2010.0, 'Subscribers per 100 People': 47.5
2	'Year': 2011.0, 'Subscribers per 100 People': 45.0	'Year': 2011.0, 'Subscribers per 100 People': 46.0
3	'Year': 2012.0, 'Subscribers per 100 People': 44.5	'Year': 2012.0, 'Subscribers per 100 People': 45.0
4	'Year': 2013.0, 'Subscribers per 100 People': 44.0	'Year': 2013.0, 'Subscribers per 100 People': 44.0

Table 11—Example for Gemini Flash predictions where it swapped the values of Turkmenistan and United States. The swapped table has score of $F_1 = 0.34$ and the corrected version has $F_1 = 0.83$.

Figure 5—Vertical Bar Chart example from PlotQA testset.

	Country	2005	2006	2007
0	Belarus	31	24	23
1	Belize	15	15	13
2	Bosnia and Herzegovina	33	34	52
3	Brazil	44	47	52
4	Cabo Verde	10	7	8
5	Canada	51	52	51

Table 12—Gemini 1.5 Flash predictions on Vertical Bar # 25905

	Country	2005 Cost of computers, communications and other services (% of comr
1	Belarus	31
2	Belize	15
3	Bosnia and Herzegovina	33
4	Brazil	45
5	Cabo Verde	11
6	Canada	52

Table 13—ChartGemma predictions for vertical bar image no. 2590 from PlotQA

Country	2005	2006	2007
Belarus	31.22	24.37	24.45
Belize	15.34	15.58	13.8
Bosnia and Herzegovina	33.22	34.26	50.3
Brazil	44.5	47.98	52.6
Cabo Verde	11.06	7.63	7.76
Canada	51.81	51.93	51.6

Table 14—Deplot predictions on Vertical Bar # 25905

	Country	Amount (in current US$)
0	Lao PDR	15000000
1	Lebanon	270000000
2	Lesotho	0
3	Liberia	0
4	Macedonia	5000000
5	Madagascar	15000000

Table 15—Gemini 1.5 Flash prediction for PlotQA vertical bar #32640

External debt owed to nonresidents by residents of different countries

Figure 6—PlotQA Vertical Bar no. 32640.

	Country	Amount (in current US$)
1	Lao PDR	100,000,000
2	Lebanon	270,000,000
3	Lesotho	0
4	Liberia	0
5	Macedonia	30,000,000
6	Madagascar	150,000,000

Table 16—ChartGemma prediction for PlotQA vertical bar #32640

Country	Debt
Lao PDR	123191000
Lebanon	2757587000
Lesotho	-11173000
Liberia	-12609000
Macedonia	33960000
Madagascar	130544000

Table 17—Deplot prediction for PlotQA vertical bar #32640

chart types. This highlights the complexity and diversity of chart-to-table tasks, which requires specialized models or very capable large general models.

Our findings have important implications for researchers and practitioners working on data extraction and analysis. They suggest that future research should focus on developing more flexible models that can handle a wide range of chart types and styles. Additionally, our results emphasize the need for larger and more diverse datasets to train and evaluate these models.

Overall, our study contributes to a deeper understanding of the challenges and opportunities in chart-to-table tasks and provides a foundation for future research in this area. We hope that our findings will fosters further innovation and improvement in data extraction and analysis.

6 LIMITATIONS

One limitation of our analysis is testing pre-trained models exclusively and did not test models like ChartOCR. These models are not directly comparable and lack existing public weights and implementations. This aspect has been deferred to future work.

7 ACKNOWLEDGMENTS

8 REFERENCES

1. Ren, Shaoqing, He, Kaiming, Girshick, Ross, and Sun, Jian (2015). "Faster r-cnn: Towards real-time object detection with region proposal networks". In: *Advances in neural information processing systems* 28.
2. Jung, Daekyoung, Kim, Wonjae, Song, Hyunjoo, Hwang, Jeong-in, Lee, Bongshin, Kim, Bohyoung, and Seo, Jinwook (2017). "Chartsense: Interactive data extraction from chart images". In: *Proceedings of the CHI Conference on Human Factors in Computing Systems*, pp. 6706–6717.
3. PACE, M (2017). "Partnership for an advanced computing environment (PACE)". In.

4. Poco, Jorge and Heer, Jeffrey (2017). "Reverse-engineering visualizations: Recovering visual encodings from chart images". In: *Computer graphics forum*. Vol. 36. 3. Wiley Online Library, pp. 353–363.

5. Vaswani, Ashish, Shazeer, Noam, Parmar, Niki, Uszkoreit, Jakob, Jones, Llion, Gomez, Aidan N, Kaiser, Łukasz, and Polosukhin, Illia (2017). "Attention is all you need". In: *Advances in neural information processing systems* 30.

6. Herzig, Jonathan, Nowak, Paweł Krzysztof, Müller, Thomas, Piccinno, Francesco, and Eisenschlos, Julian Martin (2020). "TaPas: Weakly supervised table parsing via pre-training". In: *arXiv preprint arXiv:2004.02349*.

7. Methani, Nitesh, Ganguly, Pritha, Khapra, Mitesh M, and Kumar, Pratyush (2020). "Plotqa: Reasoning over scientific plots". In: *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pp. 1527–1536.

8. Raffel, Colin, Shazeer, Noam, Roberts, Adam, Lee, Katherine, Narang, Sharan, Matena, Michael, Zhou, Yanqi, Li, Wei, and Liu, Peter J (2020). "Exploring the limits of transfer learning with a unified text-to-text transformer". In: *Journal of machine learning research* 21.140, pp. 1–67.

9. Luo, Junyu, Li, Zekun, Wang, Jinpeng, and Lin, Chin-Yew (2021). "Chartocr: Data extraction from charts images via a deep hybrid framework". In: *Proceedings of the IEEE/CVF winter conference on applications of computer vision*, pp. 1917–1925.

10. Sreevalsan-Nair, Jaya, Dadhich, Komal, and Daggubati, Siri Chandana (2021). "Tensor fields for data extraction from chart images: bar charts and scatter plots". In: *Topological Methods in Data Analysis and Visualization VI: Theory, Applications, and Software*. Springer, pp. 219–241.

11. Kumar, Anukriti, Ganu, Tanuja, and Guha, Saikat (2022). "ChartParser: Automatic Chart Parsing for Print-Impaired". In: *arXiv preprint arXiv:2211.08863*.

12. Li, Shufan, Lu, Congxi, Li, Linkai, and Zhou, Haoshuai (2022). "Chart-RCNN: Efficient Line Chart Data Extraction from Camera Images". In: *arXiv preprint arXiv:2211.14362*.

13. Liu, Fangyu, Eisenschlos, Julian Martin, Piccinno, Francesco, Krichene, Syrine, Pang, Chenxi, Lee, Kenton, Joshi, Mandar, Chen, Wenhu, Collier, Nigel, and Altun, Yasemin (2022a). "Deplot: One-shot visual language reasoning by plot-to-table translation". In: *arXiv preprint arXiv:2212.10505*.

14. Liu, Fangyu, Piccinno, Francesco, Krichene, Syrine, Pang, Chenxi, Lee, Kenton, Joshi, Mandar, Altun, Yasemin, Collier, Nigel, and Eisenschlos, Julian

Martin (2022b). "Matcha: Enhancing visual language pretraining with math reasoning and chart derendering". In: *arXiv preprint arXiv:2212.09662*.

15. Masry, Ahmed, Long, Do Xuan, Tan, Jia Qing, Joty, Shafiq, and Hoque, Enamul (2022). "Chartqa: A benchmark for question answering about charts with visual and logical reasoning". In: *arXiv preprint arXiv:2203.10244*.

16. Cheng, Zhi-Qi, Dai, Qi, and Hauptmann, Alexander G (2023). "Chartreader: A unified framework for chart derendering and comprehension without heuristic rules". In: *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 22202–22213.

17. Lee, Kenton, Joshi, Mandar, Turc, Iulia Raluca, Hu, Hexiang, Liu, Fangyu, Eisenschlos, Julian Martin, Khandelwal, Urvashi, Shaw, Peter, Chang, Ming-Wei, and Toutanova, Kristina (2023). "Pix2struct: Screenshot parsing as pretraining for visual language understanding". In: *International Conference on Machine Learning*. PMLR, pp. 18893–18912.

18. Masry, Ahmed, Kavehzadeh, Parsa, Do, Xuan Long, Hoque, Enamul, and Joty, Shafiq (2023). "Unichart: A universal vision-language pretrained model for chart comprehension and reasoning". In: *arXiv preprint arXiv:2305.14761*.

19. Rousseau, Jean-Jacques and Kapralos, Bill (2023). *Pattern Recognition, Computer Vision, and Image Processing. ICPR 2022 International Workshops and Challenges: Montreal, QC, Canada, August 21–25, 2022, Proceedings, Part I*. Vol. 13643. Springer Nature.

20. Touvron, Hugo, Martin, Louis, Stone, Kevin, Albert, Peter, Almahairi, Amjad, Babaei, Yasmine, Bashlykov, Nikolay, Batra, Soumya, Bhargava, Prajjwal, Bhosale, Shruti, et al. (2023). "Llama 2: Open foundation and fine-tuned chat models". In: *arXiv preprint arXiv:2307.09288*.

21. Voigt, Henrik, Carvalhais, Nuno, Meuschke, Monique, Reichstein, Markus, Zarrie, Sina, and Lawonn, Kai (2023). "VIST5: An adaptive, retrieval-augmented language model for visualization-oriented dialog". In: *The 2023 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, pp. 70–81.

22. Beyer, Lucas, Steiner, Andreas, Pinto, André Susano, Kolesnikov, Alexander, Wang, Xiao, Salz, Daniel, Neumann, Maxim, Alabdulmohsin, Ibrahim, Tschannen, Michael, Bugliarello, Emanuele, et al. (2024). "Paligemma: A versatile 3b vlm for transfer". In: *arXiv preprint arXiv:2407.07726*.

23.   Cao, Yukun, Han, Shuo, Gao, Zengyi, Ding, Zezhong, Xie, Xike, and Zhou, S Kevin (2024). "Graphinsight: Unlocking insights in large language models for graph structure understanding". In: *arXiv preprint arXiv:2409.03258*.

24.   Chung, Hyung Won, Hou, Le, Longpre, Shayne, Zoph, Barret, Tay, Yi, Fedus, William, Li, Yunxuan, Wang, Xuezhi, Dehghani, Mostafa, Brahma, Siddhartha, et al. (2024). "Scaling instruction-finetuned language models". In: *Journal of Machine Learning Research* 25.70, pp. 1–53.

25.   Masry, Ahmed, Shahmohammadi, Mehrad, Parvez, Md Rizwan, Hoque, Enamul, and Joty, Shafiq (2024a). "ChartInstruct: Instruction Tuning for Chart Comprehension and Reasoning". In: *arXiv preprint arXiv:2403.09028*.

26.   Masry, Ahmed, Thakkar, Megh, Bajaj, Aayush, Kartha, Aaryaman, Hoque, Enamul, and Joty, Shafiq (2024b). "ChartGemma: Visual Instruction-tuning for Chart Reasoning in the Wild". In: *arXiv preprint arXiv:2407.04172*.

27.   Team, Gemini, Georgiev, Petko, Lei, Ving Ian, Burnell, Ryan, Bai, Libin, Gulati, Anmol, Tanzer, Garrett, Vincent, Damien, Pan, Zhufeng, Wang, Shibo, et al. (2024). "Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context". In: *arXiv preprint arXiv:2403.05530*.

28.   Goswami, Kanika, Mathur, Puneet, Rossi, Ryan, and Dernoncourt, Franck (2025). "ChartCitor: Multi-Agent Framework for Fine-Grained Chart Visual Attribution". In: *arXiv preprint arXiv:2502.00989*.

29.   Zhao, Xuanle, Luo, Xianzhen, Shi, Qi, Chen, Chi, Wang, Shuo, Che, Wanxiang, Liu, Zhiyuan, and Sun, Maosong (2025). "ChartCoder: Advancing Multimodal Large Language Model for Chart-to-Code Generation". In: *arXiv preprint arXiv:2501.06598*.

A   EVALUATION METRICS

A.1   Relative Mapping Similarity (RMS)

A.2   metrics introduced in ChartOCR

In this subsection, we illustrate metrics introduced by Luo et al., 2021.

A.2.1   *Bar Chart Metric*

Given a predicted bounding box $p = [x_p, y_p, w_p, h_p]$ and a ground truth box $g = [x_g, y_g, w_g, h_g]$, the distance function $D(p, g)$ is defined as:

Require: prp (Predicted Table), t (Ground Truth Table)

Ensure: $rms\_precision, rms\_recall, rms\_f1$

Compute Normalized Edit Distance between $p_r \| p_c$ and $t_r \| t_c$, where $\|$ is concatenation operator. Compute pairwise similarities matrix make binarized similarities matrix by inserting 1 in place of highest similarities and zeros otherwise.

1: **for do** p_i in p.values:

2:    **for do** t_j in t.values: $d_{theta_{ij}} = \min(1, \frac{|p_i - t_j|}{|t_j|})$ $d_{tao_{theta_{ij}}} = (1 - NormalizedEditDistance) * (1 - d_{theta})$

3:    **end for**

4: **end for** $RMS_{precision} = \sum_i \sum_j d_{tao_{theta_{ij}}}/len(p)$ $RMS_{recall} = \sum_i \sum_j d_{tao_{theta_{ij}}}/len(t)$ $RMS_{f1} = 2 * \frac{RMS_{precision} * RMS_{recall}}{RMS_{precision} + RMS_{recall}}$

$$D(p, g) = \min\left(1, \left|\frac{x_p - x_g}{w_g}\right| + \left|\frac{y_p - y_g}{h_g}\right| + \left|\frac{h_p - h_g}{h_g}\right|\right)$$

Width w is ignored because it does not affect chart reading.

Construct a pairwise cost matrix C between predicted and ground truth boxes, then solve a minimum cost assignment problem using the Hungarian algorithm. The final score is:

$$Score = 1 - \frac{\min cost}{K}, \quad \text{where } K = \max(N, M)$$

Here N and M are the number of predicted and ground truth boxes respectively.

A.2.2 *Line Chart Metric*

Let $P = \{(x_1, y_1), \ldots, (x_N, y_N)\}$ be the predicted points and $G = \{(u_1, v_1), \ldots, (u_M, v_M)\}$ be the ground truth. The recall is defined as:

$$Rec(P, G) = \frac{1}{u_M - u_1} \sum_{i=1}^{M} (1 - Err(v_i, u_i, P)) \cdot Intv(i, G)$$

Where:

$$\text{Err}(v_i, u_i, P) = \min\left(1, \left|\frac{v_i - I(P, u_i)}{v_i}\right|\right)$$

$I(P, u_i)$ is the interpolated value of P at $x = u_i$. $\text{Intv}(i, G)$ is the interval weight:

$$\text{Intv}(i, G) = \begin{cases} \frac{u_{i+1} - u_i}{2} & \text{if } i = 1 \\ \frac{u_i - u_{i-1}}{2} & \text{if } i = M \\ \frac{u_{i+1} - u_{i-1}}{2} & \text{otherwise} \end{cases}$$

Precision is defined similarly. The final score is the F1-score:

$$\text{F1} = \frac{2 \cdot \text{Prec} \cdot \text{Rec}}{\text{Prec} + \text{Rec}}$$

A.2.3 *Pie Chart Metric*

Let $P = [x_1, \ldots, x_N]$ and $G = [y_1, \ldots, y_M]$ be the predicted and ground-truth value sequences in clockwise order. The match score is computed via dynamic programming:

$$\text{score}(i, j) = \max\left(\text{score}(i{-}1, j), \text{score}(i, j{-}1), \text{score}(i{-}1, j{-}1) + 1 - \left|\frac{x_i - y_j}{y_j}\right|\right)$$

with $\text{score}(0, j) = \text{score}(i, 0) = 0$ for all i, j. The final normalized matching score is:

$$\text{Score} = \frac{\text{score}(N, M)}{M}$$

B EXAMPLES FOR FAILING PLOT TYPES

This section includes examples for all chart types that can not be processed by any of tested models.

B.1 Scatter Plots

Ground Truth csv:

```
x,y
BS71\n (3),1.9980978260869584
```

Figure 7—Scatter Plot from ICPR22 testset #2279

DePlot output csv:

```
X,Y,Label
2,1,"1
 8871
(3)"
2,1,"3
 8874
(4)"
3,1,"2
 8873
(6)"
2,1,"1
 8874
(6)"
2,1,"1
 8868
(9)"
2,1,"1
 8868
(9)"
2,1,"1
 88290
(9)"
2,1,"1
 88291
(9)"
2,1,"1
 88434
(9)"
2,1,"1
 88437
(9)"
2,1,"1
 8870
(11)"
```

```
2,1,"1
 8860
(14)"
1,1,"1
 8875
(19)"
1,1,"1
 8872
(23)"
1,1,"1
 88293
(NA)"
5,1,1
```

Gemini 1.5 Flash output json:

```
[{"Strain": "BS71", "Serotype": 3, "Days to Onset": 2, "Color": "Blue"}, {"Strain": "TI
```

ChartGemma output text:

```
The table should have two columns: 'Strain' and 'Days to Osse'. Each row in the table s
```

B.2 Box Plots

1. Gemini 1.5 Flash was the only model to correctly extract data points from vertical box plots. However, comparing Gemini 1.5 Flash's output to ground truth table needs custom logic because some graphs have more than one data series.

Ground Truth csv:

```
series_name,subplot_name,first_quartile,max,median,min,third_quartile
[unnamed data series #0],Non - carriers,0.11198945981554975,15.131752305665346,5.961791
[unnamed data series #0],carriers,7.964426877470356,21.034255599472985,17.8722002635046
```

DePlot prediction:

```
Metric,Group,Value
```

Figure 8—Vertical Box Plot from ICPR22 testset #1855

```
"Salt sensitivity in mm<0xE2><0x80><0xA0>
Median(IQR)
 Non- carriers
(n=34)",carriers (n=5),10.0
"Salt sensitivity in mm<0xE2><0x80><0xA0>
Median(IQR)
 Non- carriers
(n=34)",Group1,21.0
```

Gemini 1.5 Flash prediction:

```
[{"Group": "Non-carriers", "Minimum": -8.5, "First Quartile (Q1)": 2, "Median": 6, "Thi
```

ChartGemma's output markdown:

```
The table should have two columns: 'Non-carriers' and 'Carriers'. Each row in the table
```